

自我來黃州已過三寒  
食年、欲惜春、意不  
容惜今年又苦雨、月社  
簫瑟、河海、棠花泥  
污、遊支雪、閣中偷負  
多夜半、真有力、何殊少  
年、病起、頭白  
春江欲入户、雨勢未  
止、雨小屋如漚、舟濫  
水雲裏、空庭裏、寒葉  
破、竈燒酒、華那  
知是寒食、但見烏  
銜、帛、天門深  
九重、噴、蒼生、在、萬里、遙  
哭、淪、窮、所、不、吹、不  
起

右黃州寒食二首

# 信息检索

## Information Retrieval

教师：孙茂松

Tel: 62781286

Email: sms@tsinghua.edu.cn

TA：胡锦涛

Email: hu-jy21@mails.tsinghua.edu.cn

# 郑重声明

- 此课件仅供选修清华大学计算机系本科生课《信息检索》的学生个人学习使用，所以只允许学生将其下载、存贮在自己的电脑中。未经孙茂松本人同意，任何人不得以任何方式扩散之（包括放到9#服务器上）。否则，由此可能引起的一切涉及知识产权的法律责任，概由该人负责。
- 此课件仅限孙茂松本人讲课使用。除孙茂松本人外，凡授课过程中，PTT文件显示此《郑重声明》之情形，即为侵权使用。



# 第三章 文本分析及自动标引 (Part 4)

## 3.5 Thesaurus及term自动关联(Word2Vec)

### Representing words as discrete symbols

In traditional NLP, we regard words as discrete symbols:  
*hotel*, *conference*, *motel* – a *localist* representation

Means one 1, the rest 0s

Such symbols for words can be represented by *one-hot* vectors:

*motel* = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]  
*hotel* = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]

Vector dimension = number of words in vocabulary (e.g., 500,000)

# 3.5 Thesaurus及term自动关联(Word2Vec)

## Problem with words as discrete symbols

**Example:** in web search, if user searches for “Seattle motel”, we would like to match documents containing “Seattle hotel”

But:

motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]

hotel = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]

These two vectors are orthogonal

There is no natural notion of **similarity** for one-hot vectors!

### Solution:

- Could try to rely on WordNet’s list of synonyms to get similarity?
  - But it is well-known to fail badly: incompleteness, etc.
- **Instead: learn to encode similarity in the vectors themselves**

# 3.5 Thesaurus及term自动关联(Word2Vec)

## Representing words by their context



- **Distributional semantics:** A word's meaning is given by the words that frequently appear close-by
  - “You shall know a word by the company it keeps” (J. R. Firth 1957: 11)
  - One of the most successful ideas of modern statistical NLP!
- When a word  $w$  appears in a text, its **context** is the set of words that appear nearby (within a fixed-size window).
- Use the many contexts of  $w$  to build up a representation of  $w$

...government debt problems turning into **banking** crises as happened in 2009...  
...saying that Europe needs unified **banking** regulation to replace the hodgepodge...  
...India has just given its **banking** system a shot in the arm...

These **context words** will represent **banking**



## 3.5 Thesaurus及term自动关联(Word2Vec)

### Word vectors

We will build a dense vector for each word, chosen so that it is similar to vectors of words that appear in similar contexts

*banking* =

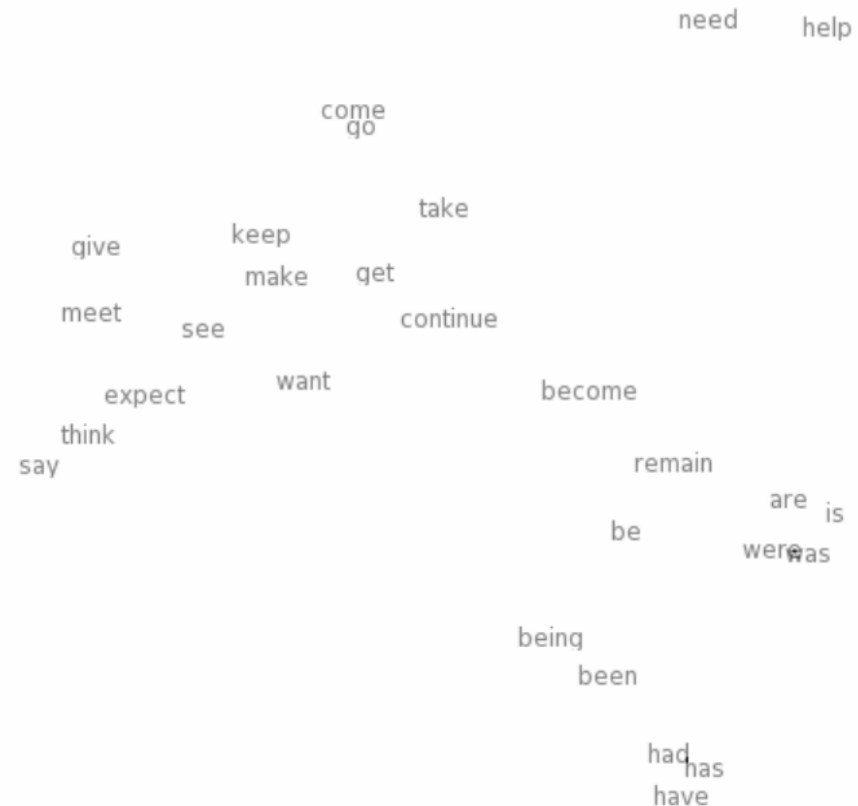
$$\begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

Note: **word vectors** are also called **word embeddings** or **(neural) word representations**  
They are a **distributed** representation

# 3.5 Thesaurus及term自动关联(Word2Vec)

## Word meaning as a neural word vector – visualization

*expect* =

$$\begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \\ 0.487 \end{pmatrix}$$




# 3.5 Thesaurus及term自动关联(Word2Vec)

Word2vec (Mikolov et al. 2013) is a framework for learning word vectors

---

## Efficient Estimation of Word Representations in Vector Space

---

**Tomas Mikolov**

Google Inc., Mountain View, CA  
tmikolov@google.com

**Kai Chen**

Google Inc., Mountain View, CA  
kaichen@google.com

**Greg Corrado**

Google Inc., Mountain View, CA  
gcorrado@google.com

**Jeffrey Dean**

Google Inc., Mountain View, CA  
jeff@google.com

---

## Distributed Representations of Words and Phrases and their Compositionality

---

**Tomas Mikolov**

Google Inc.  
Mountain View  
mikolov@google.com

**Ilya Sutskever**

Google Inc.  
Mountain View  
ilyasu@google.com

**Kai Chen**

Google Inc.  
Mountain View  
kai@google.com

**Greg Corrado**

Google Inc.  
Mountain View  
gcorrado@google.com

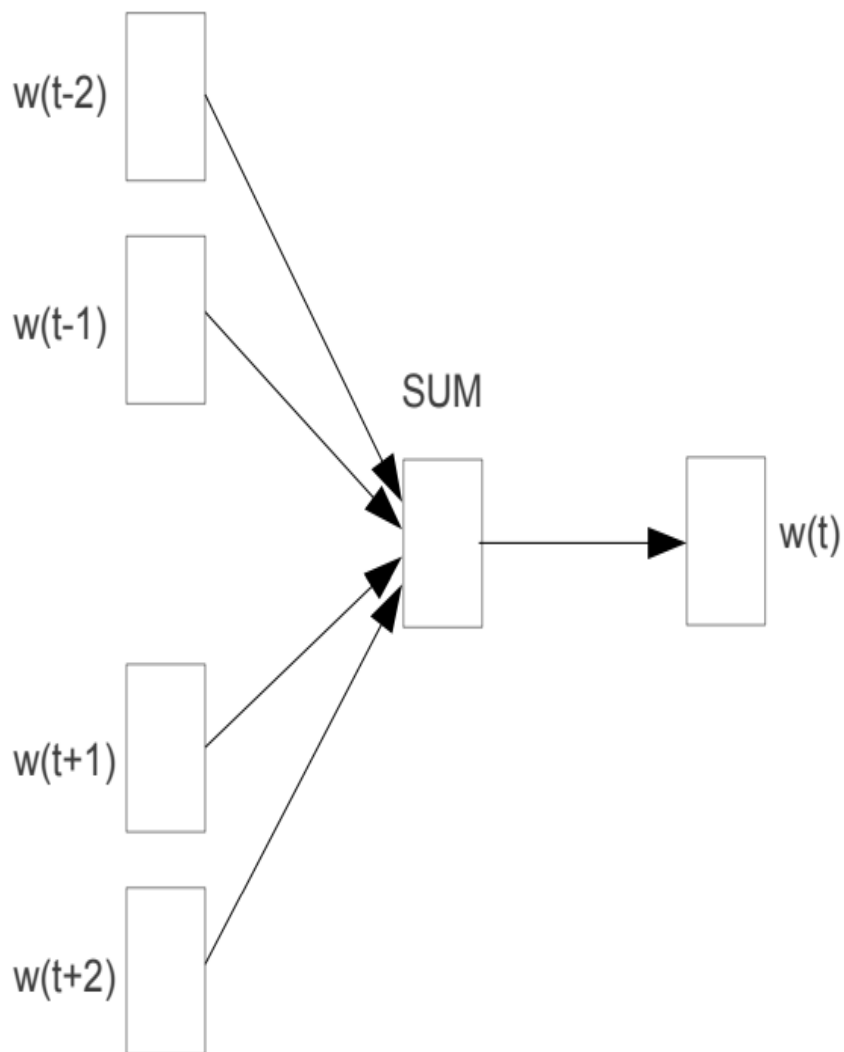
**Jeffrey Dean**

Google Inc.  
Mountain View  
jeff@google.com

INPUT

PROJECTION

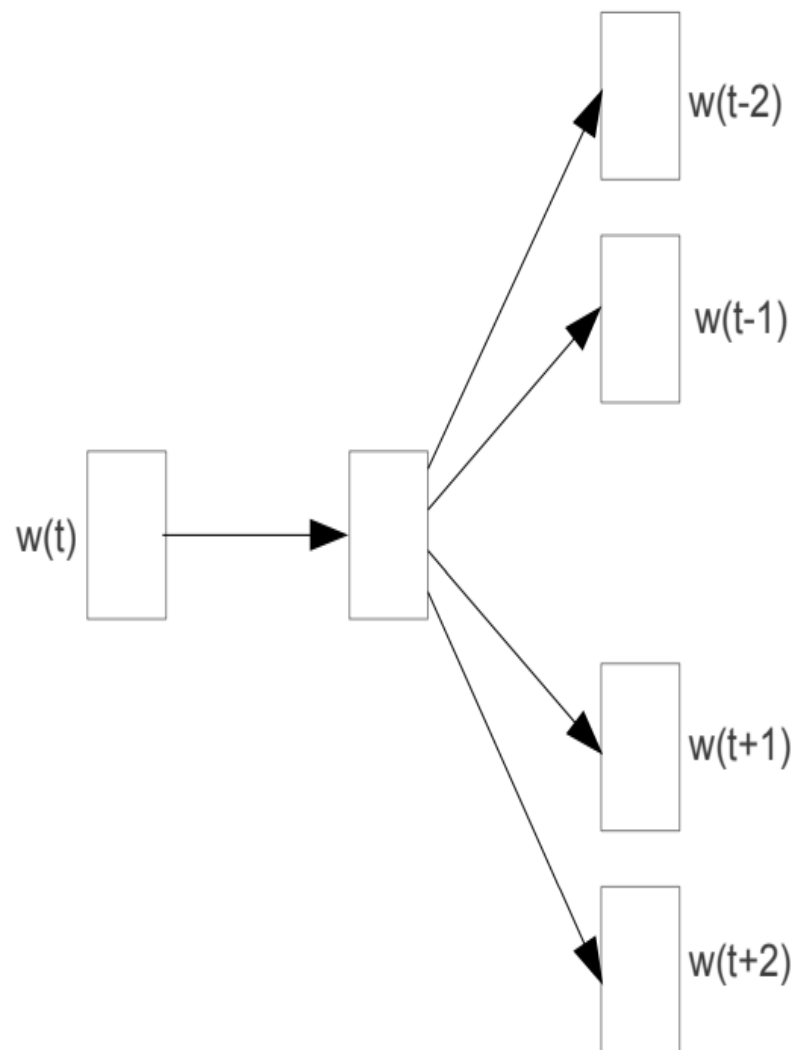
OUTPUT

**CBOW**

INPUT

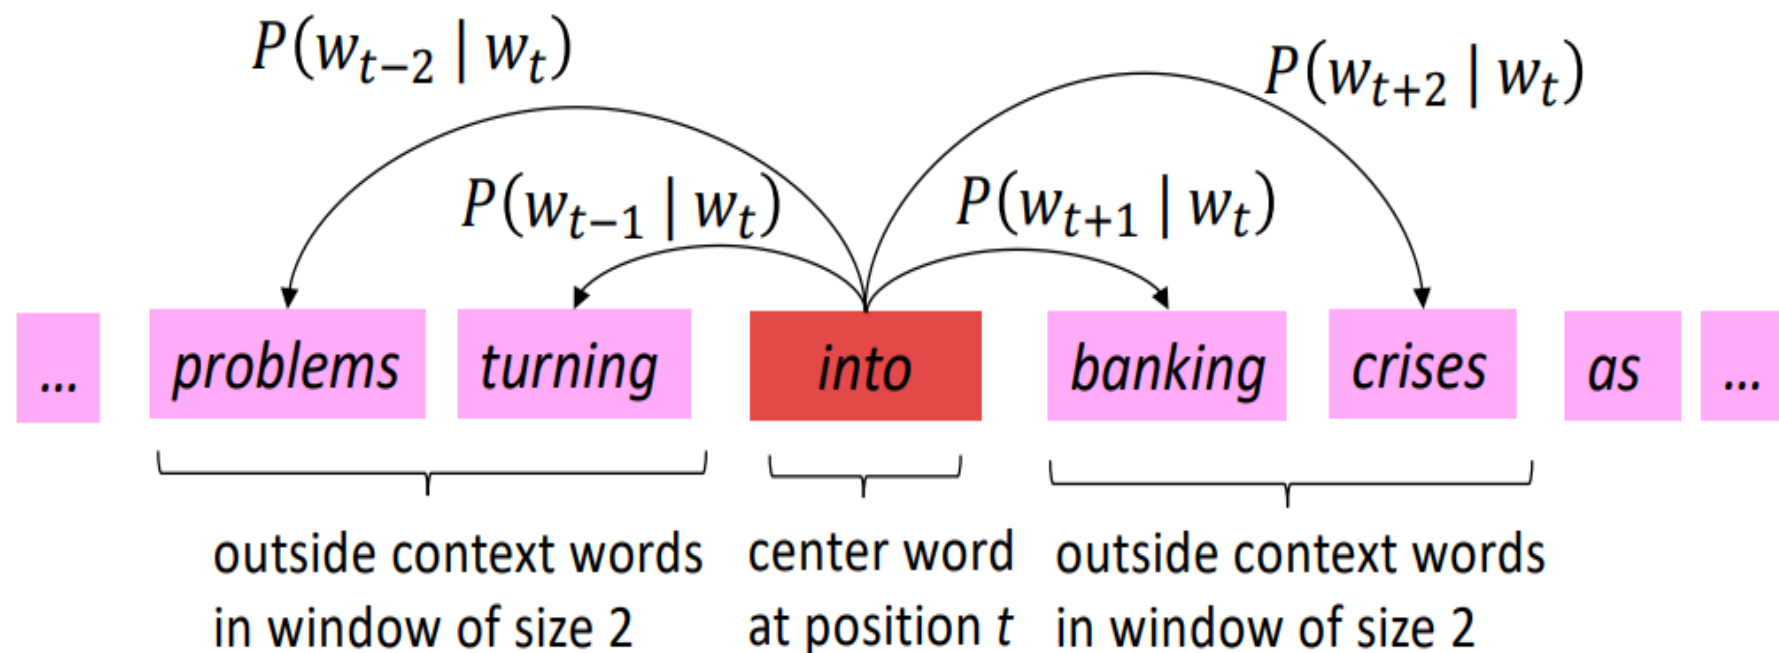
PROJECTION

OUTPUT

**Skip-gram**

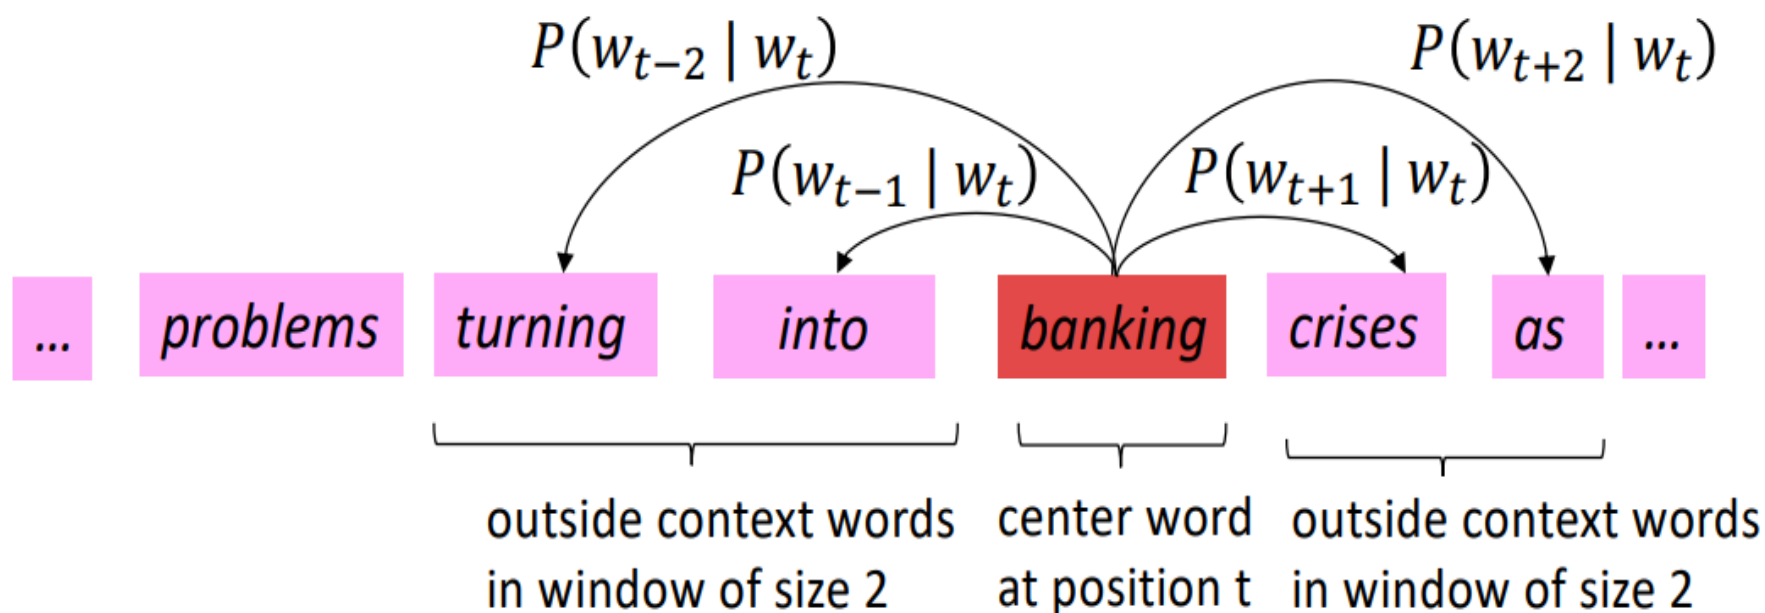
## 3.5 Thesaurus及term自动关联(Word2Vec)

Example windows and process for computing  $P(w_{t+j} | w_t)$



# 3.5 Thesaurus及term自动关联(Word2Vec)

Example windows and process for computing  $P(w_{t+j} | w_t)$



# Word2vec: objective function

For each position  $t = 1, \dots, T$ , predict context words within a window of fixed size  $m$ , given center word  $w_j$ . Data likelihood:

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m \\ j \neq 0}} P(w_{t+j} | w_t; \theta)$$

$\theta$  is all variables  
to be optimized

sometimes called a *cost* or *loss* function

The **objective function**  $J(\theta)$  is the (average) negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

Minimizing objective function  $\Leftrightarrow$  Maximizing predictive accuracy

# Word2vec: objective function

- We want to minimize the objective function:

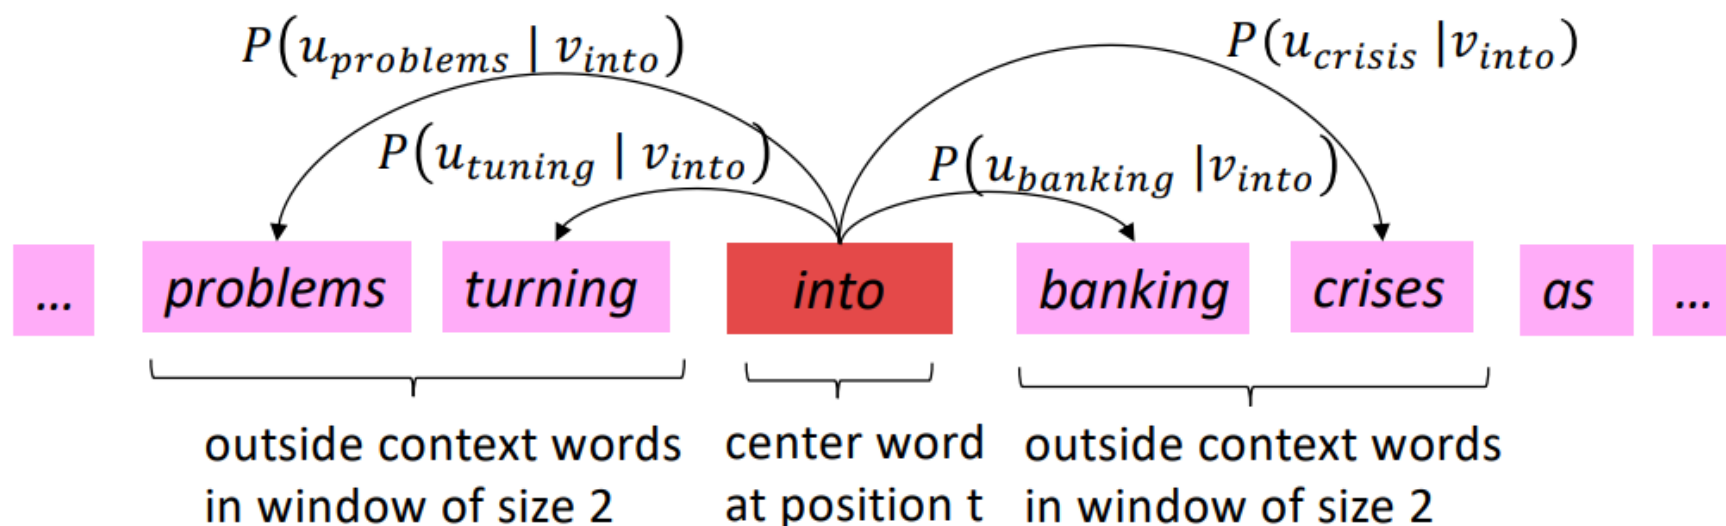
$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

- **Question:** How to calculate  $P(w_{t+j} | w_t; \theta)$  ?
- **Answer:** We will *use two* vectors per word  $w$ :
  - $v_w$  when  $w$  is a center word
  - $u_w$  when  $w$  is a context word
- Then for a center word  $c$  and a context word  $o$ :

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

## 3.5 Thesaurus及term自动关联(Word2Vec)

- Example windows and process for computing  $P(w_{t+j} | w_t)$
- $P(u_{problems} | v_{into})$  short for  $P(problems | into ; u_{problems}, v_{into}, \theta)$





# 3.5 Thesaurus及term自动关联(Word2Vec)

## Word2vec: prediction function

② Exponentiation makes anything positive

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

① Dot product compares similarity of  $o$  and  $c$ .

$$u^T v = u \cdot v = \sum_{i=1}^n u_i v_i$$

Larger dot product = larger probability

③ Normalize over entire vocabulary to give probability distribution

- This is an example of the **softmax function**  $\mathbb{R}^n \rightarrow (0,1)^n$  Open region

$$\text{softmax}(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)} = p_i$$

- The softmax function maps arbitrary values  $x_i$  to a probability distribution  $p_i$ 
  - “max” because amplifies probability of largest  $x_i$
  - “soft” because still assigns some probability to smaller  $x_i$
  - Frequently used in Deep Learning

But sort of a weird name because it returns a distribution!

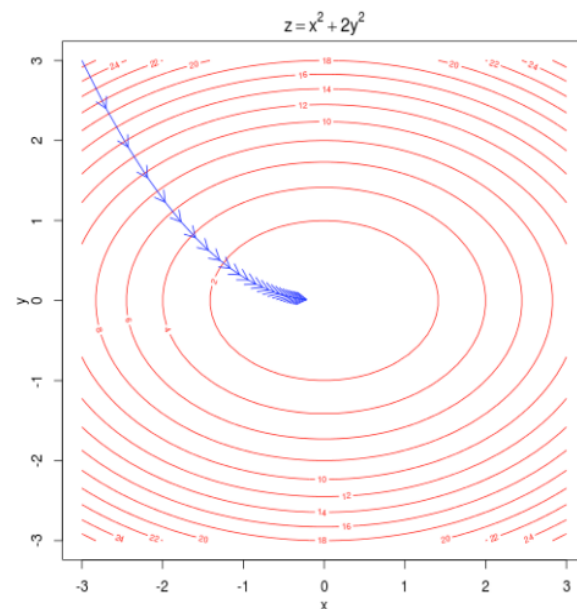
# 3.5 Thesaurus及term自动关联(Word2Vec)

To train the model: Optimize value of parameters to minimize loss

To train a model, we gradually adjust parameters to minimize a loss

- Recall:  $\theta$  represents **all** the model parameters, in one long vector
- In our case, with  $d$ -dimensional vectors and  $V$ -many words, we have:
- Remember: every word has two vectors

$$\theta = \begin{bmatrix} v_{aardvark} \\ v_a \\ \vdots \\ v_{zebra} \\ u_{aardvark} \\ u_a \\ \vdots \\ u_{zebra} \end{bmatrix} \in \mathbb{R}^{2dV}$$



- We optimize these parameters by walking down the gradient (see right figure)
- We compute **all** vector gradients!

# Chain Rule

- Chain rule! If  $y = f(u)$  and  $u = g(x)$ , i.e.,  $y = f(g(x))$ , then:

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx} = \frac{df(u)}{du} \frac{dg(x)}{dx}$$

- Simple example:  $\frac{dy}{dx} = \frac{d}{dx} 5(x^3 + 7)^4$

$$y = f(u) = 5u^4$$

$$u = g(x) = x^3 + 7$$

$$\frac{dy}{du} = 20u^3$$

$$\frac{du}{dx} = 3x^2$$

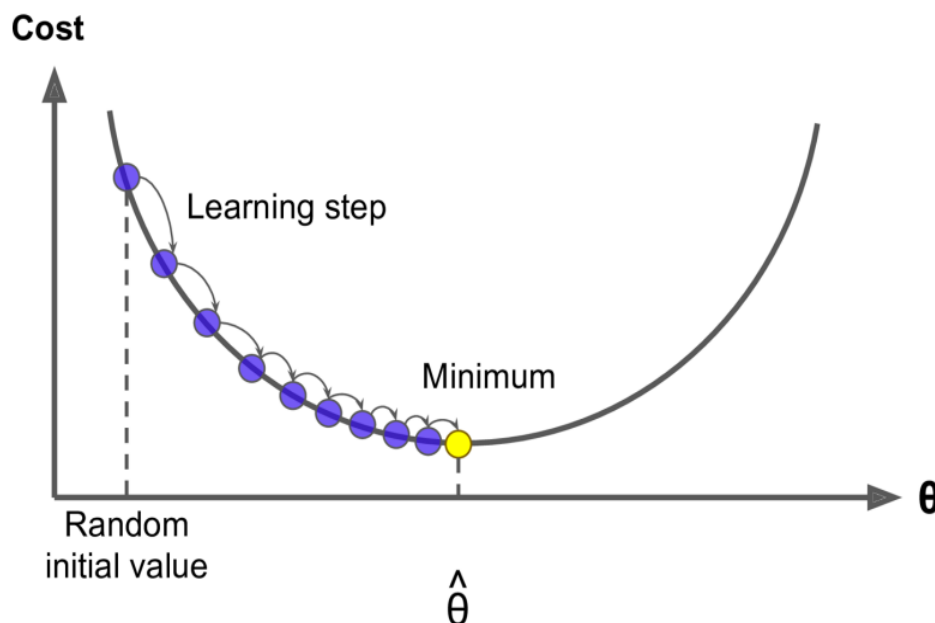
$$\frac{dy}{dx} = 20(x^3 + 7)^3 \cdot 3x^2$$

Useful basic fact:  $\frac{\partial \mathbf{x}^T \mathbf{a}}{\partial \mathbf{x}} = \frac{\partial \mathbf{a}^T \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$

# 3.5 Thesaurus及term自动关联(Word2Vec)

## Gradient Descent

- We have a cost function  $J(\theta)$  we want to minimize
- **Gradient Descent** is an algorithm to minimize  $J(\theta)$
- **Idea:** for current value of  $\theta$ , calculate gradient of  $J(\theta)$ , then take **small step in direction of negative gradient**. Repeat.



Note: Our objectives may not be convex like this ☹️

But life turns out to be okay 😊

## 3.5 Thesaurus及term自动关联(Word2Vec)

### Gradient Descent

- Update equation (in matrix notation):

$$\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$$

$\alpha$  = *step size* or *learning rate*

- Update equation (for single parameter):

$$\theta_j^{new} = \theta_j^{old} - \alpha \frac{\partial}{\partial \theta_j^{old}} J(\theta)$$

## 3.5 Thesaurus及term自动关联(Word2Vec)

- Iteratively take gradients at each such window for SGD
- But in each window, we only have at most  $2m + 1$  words, so  $\nabla_{\theta} J_t(\theta)$  is very sparse!

$$\nabla_{\theta} J_t(\theta) = \begin{bmatrix} 0 \\ \vdots \\ \nabla_{v_{like}} \\ \vdots \\ 0 \\ \nabla_{u_I} \\ \vdots \\ \nabla_{u_{learning}} \\ \vdots \end{bmatrix} \in \mathbb{R}^{2dV}$$

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

则对其中两个词 $\mathbf{o}$ 和 $\mathbf{c}$ , 令:

$$P(O = \mathbf{o} | C = \mathbf{c}) = \frac{\exp(u_{\mathbf{o}}^T v_{\mathbf{c}})}{\sum_{x=1}^V \exp(u_x^T v_{\mathbf{c}})}$$

现计算其导数, 即:

$$\frac{\partial}{\partial v_{\mathbf{c}}} \log \frac{\exp(u_{\mathbf{o}}^T v_{\mathbf{c}})}{\sum_{x=1}^V \exp(u_x^T v_{\mathbf{c}})}$$

注意: 向量右上角 $T$ 表示转置,  $J$ 中 $T$ 表示语料库规模。下同



可分解为两部分。其中前半部分：

$$\frac{\partial}{\partial v_c} (\log (\exp(u_o^T v_c))) = \frac{\partial}{\partial v_c} (u_o^T v_c) = u_o$$

其中后半部分：

$$\begin{aligned} & - \frac{1}{\sum_{x=1}^V \exp(u_x^T v_c)} \sum_{x=1}^V \frac{\partial}{\partial v_c} \exp(u_x^T v_c) \\ &= - \sum_{x=1}^V \frac{\exp(u_x^T v_c)}{\sum_{x=1}^V \exp(u_x^T v_c)} * u_x \\ &= - \sum_{x=1}^V P(x|c) * u_x \end{aligned}$$

于是有：

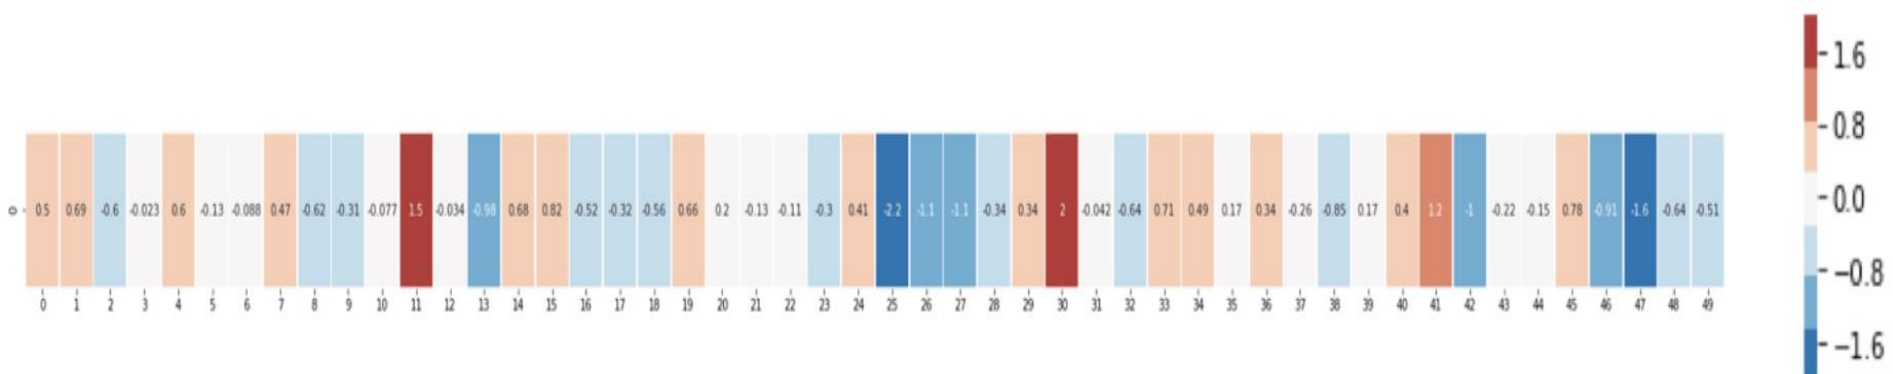
$$\frac{\partial}{\partial v_c} J_t(\theta) = - \sum_{o \text{ 在 } c \text{ 窗口内}} \left( u_o - \sum_{x=1}^V P(x|c) * u_x \right)$$

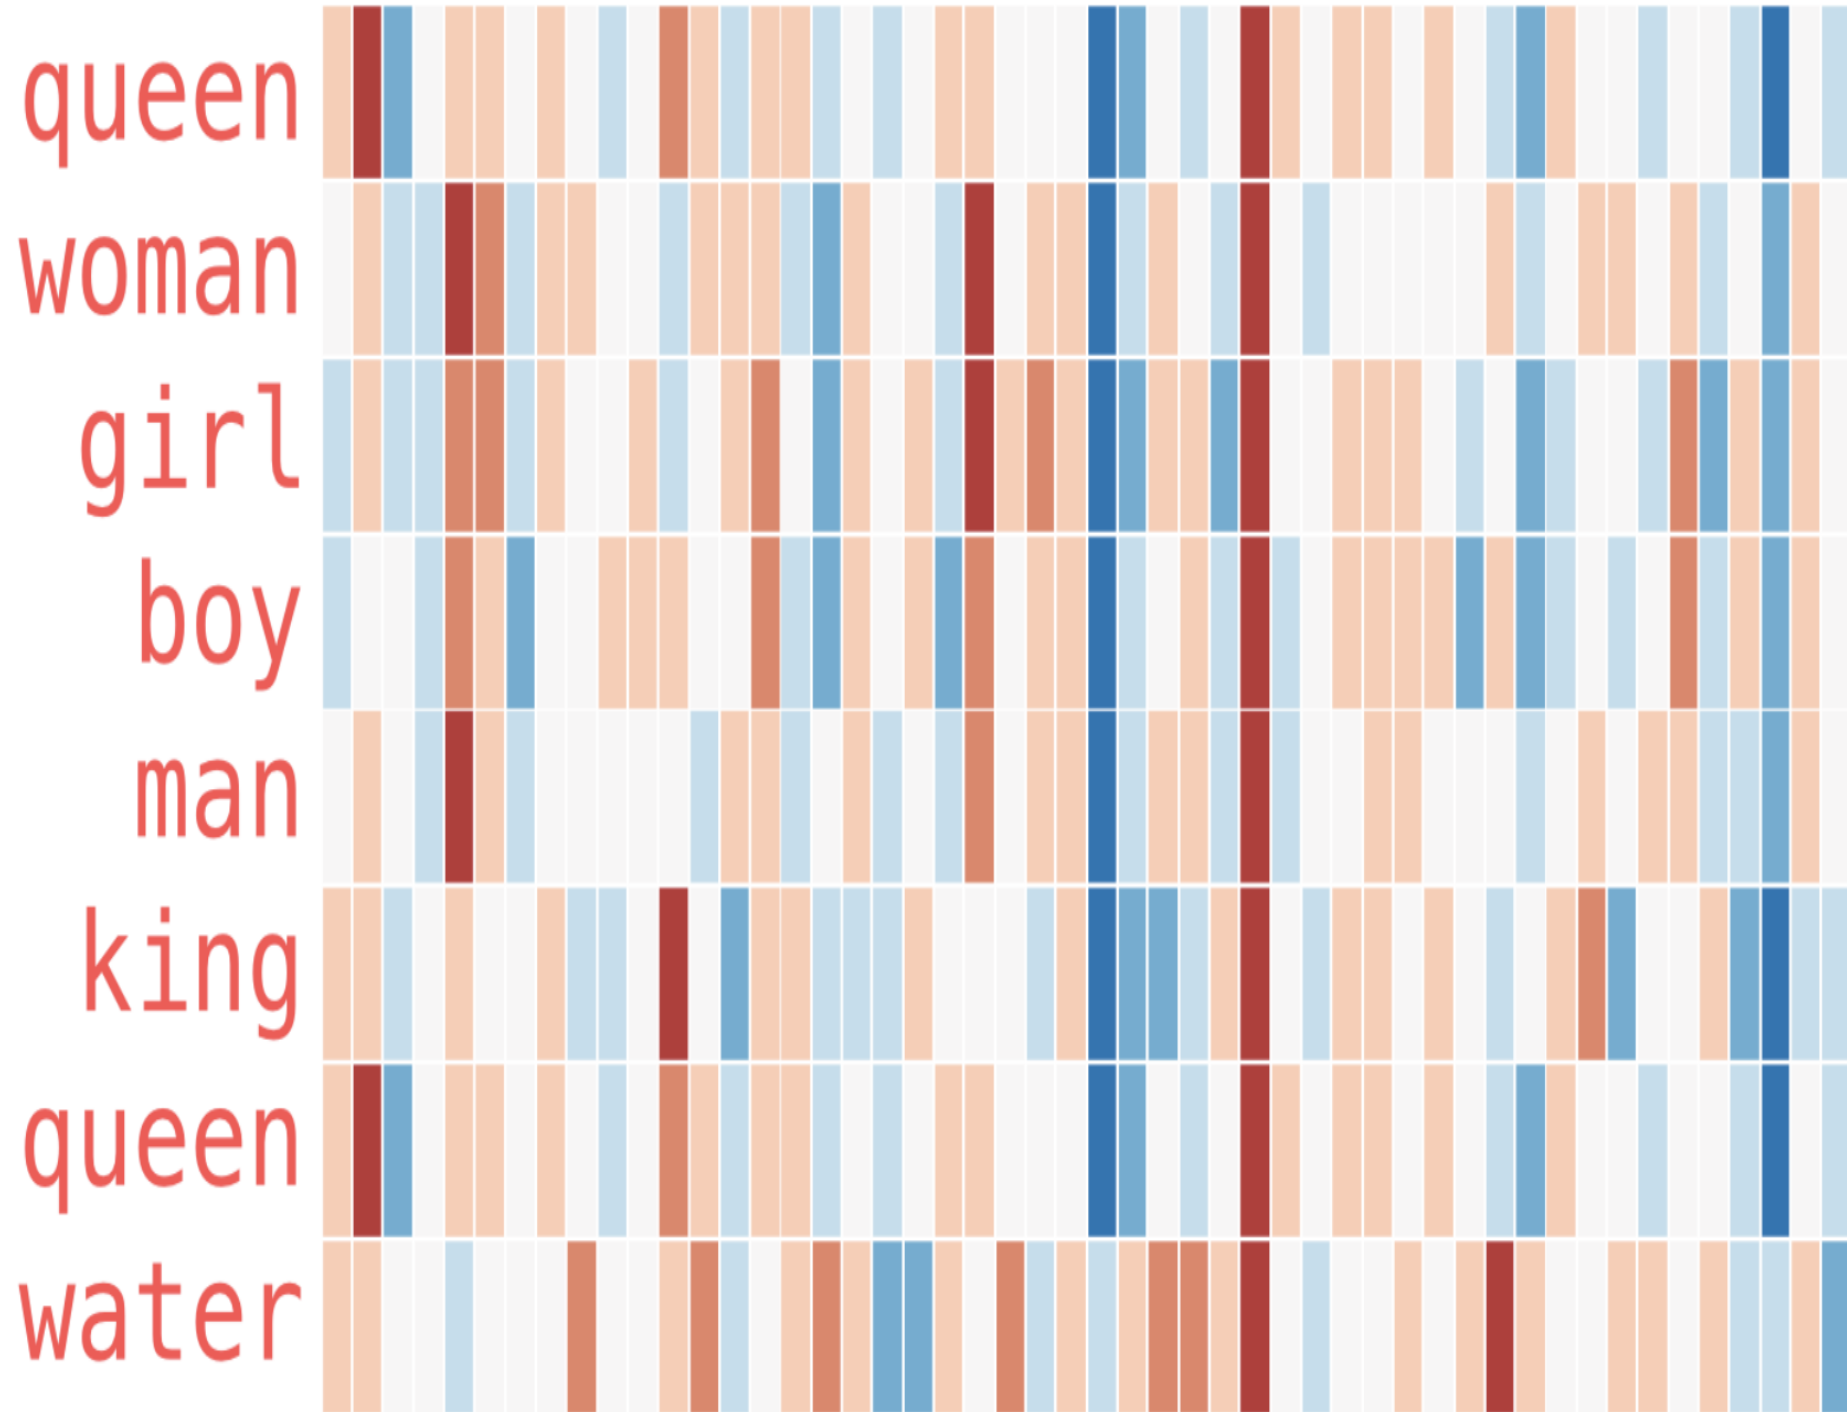
类似地，可得：

$$\frac{\partial}{\partial u_o} J_t(\theta) = - \sum_{o \text{ 在 } c \text{ 窗口内}} (v_c - P(x|c) * v_c)$$

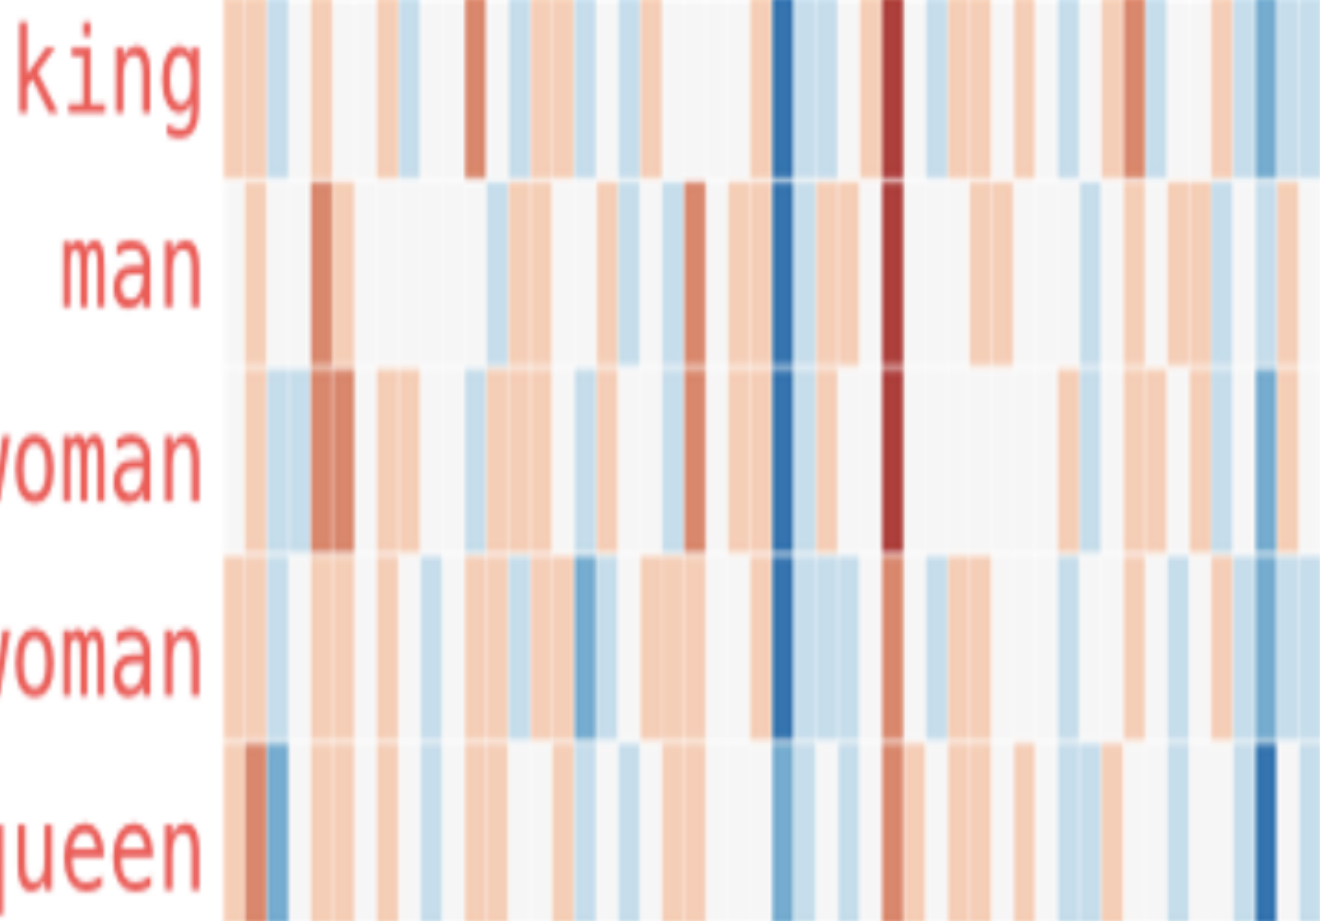
This is a word embedding for the word “king” (GloVe vector trained on Wikipedia):

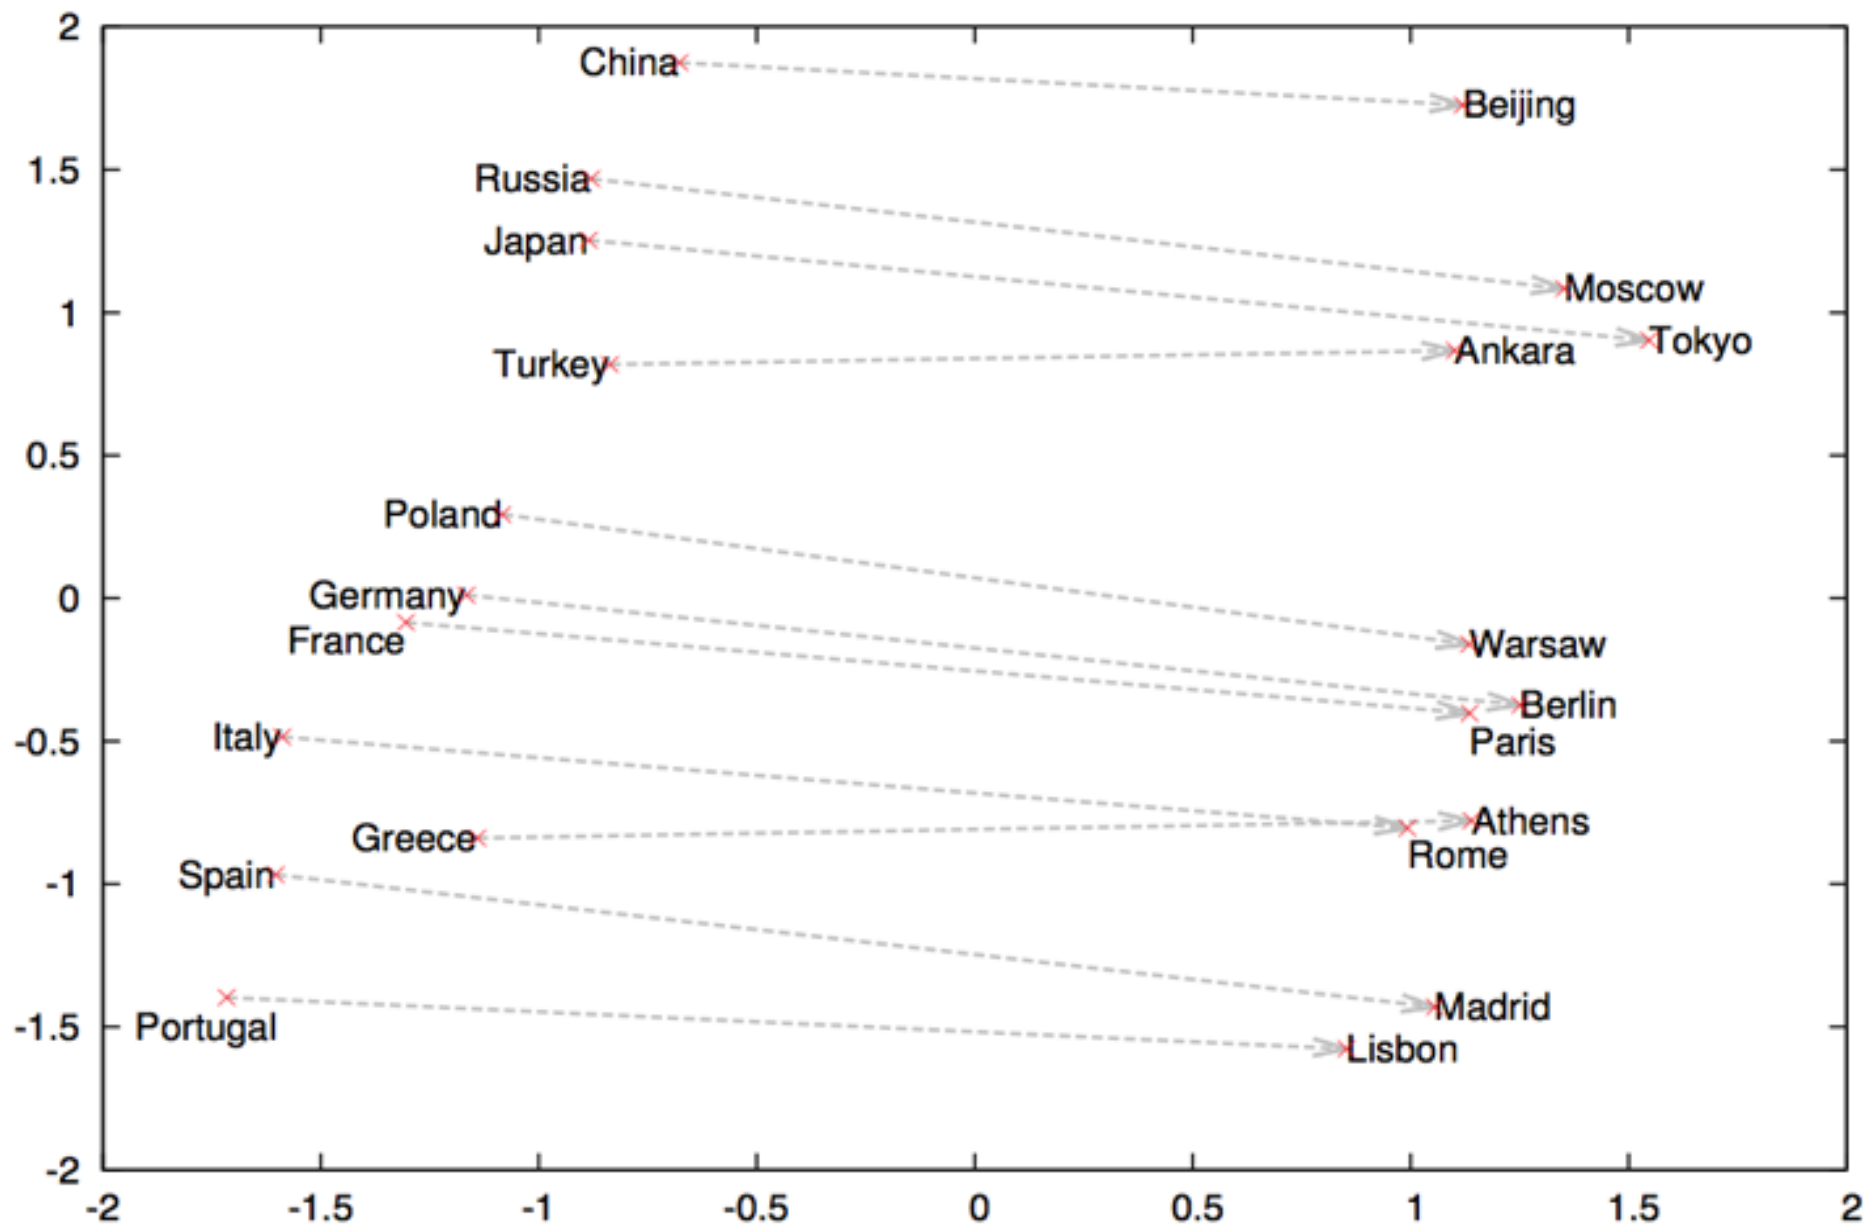
```
[ 0.50451 , 0.68607 , -0.59517 , -0.022801, 0.60046 , -0.13498 , -0.08813 , 0.47377 , -0.61798 , -0.31012 ,  
-0.076666, 1.493 , -0.034189, -0.98173 , 0.68229 , 0.81722 , -0.51874 , -0.31503 , -0.55809 , 0.66421 , 0.1961  
, -0.13495 , -0.11476 , -0.30344 , 0.41177 , -2.223 , -1.0756 , -1.0783 , -0.34354 , 0.33505 , 1.9927 ,  
-0.04234 , -0.64319 , 0.71125 , 0.49159 , 0.16754 , 0.34344 , -0.25663 , -0.8523 , 0.1661 , 0.40102 , 1.1685 ,  
-1.0137 , -0.21585 , -0.15155 , 0.78321 , -0.91241 , -1.6106 , -0.64426 , -0.51042 ]
```





king - man + woman  $\approx$  queen





$$W(\text{“China”}) - W(\text{“Beijing”}) \simeq W(\text{“Japan”}) - W(\text{“Tokyo”})$$

# Intrinsic word vector evaluation

- Word Vector Analogies

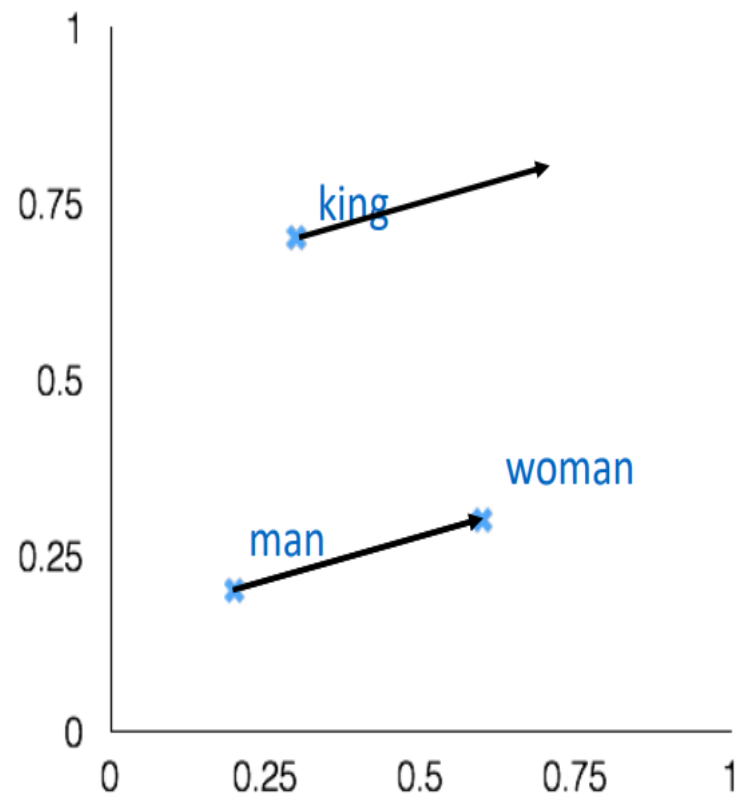
a:b :: c:?



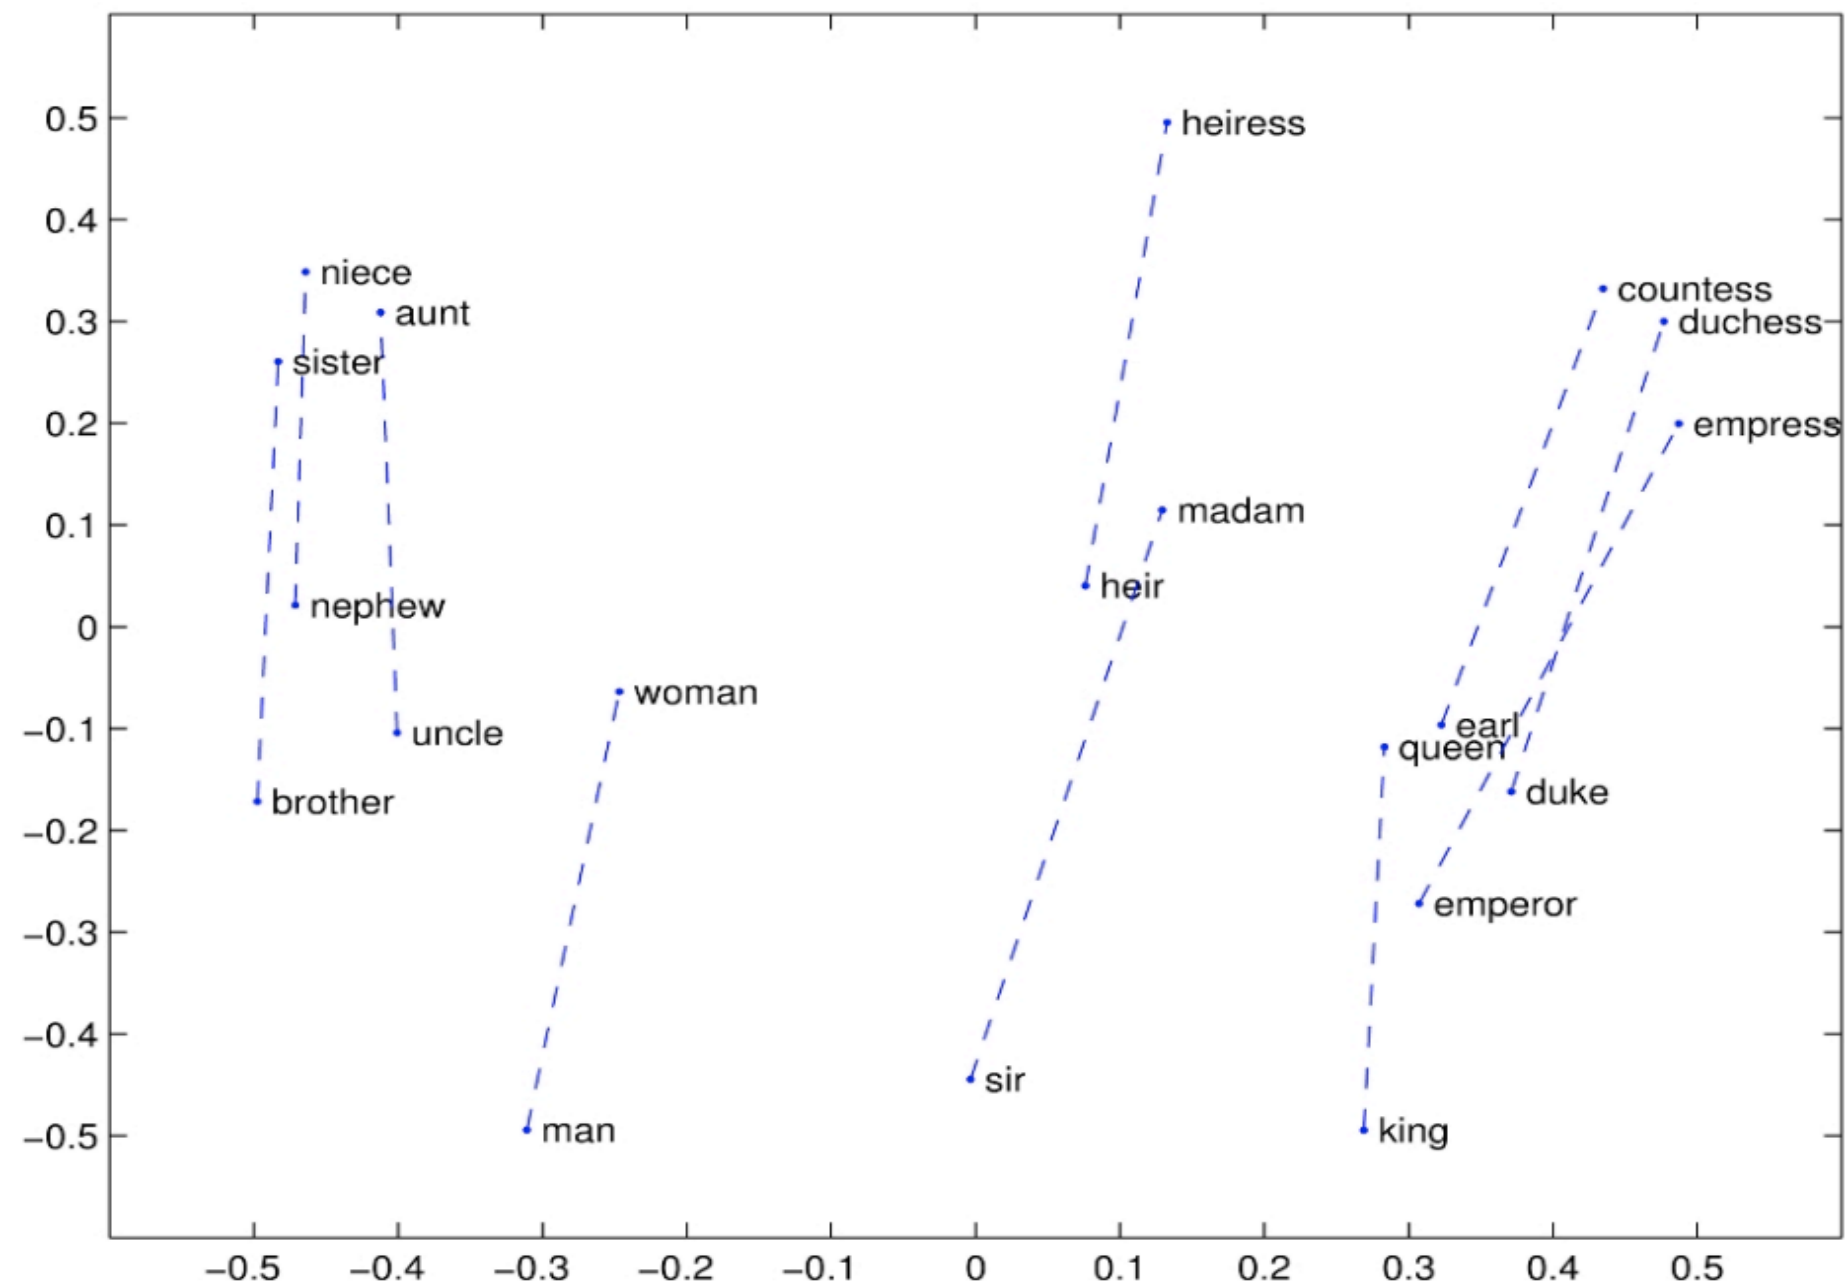
$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}$$

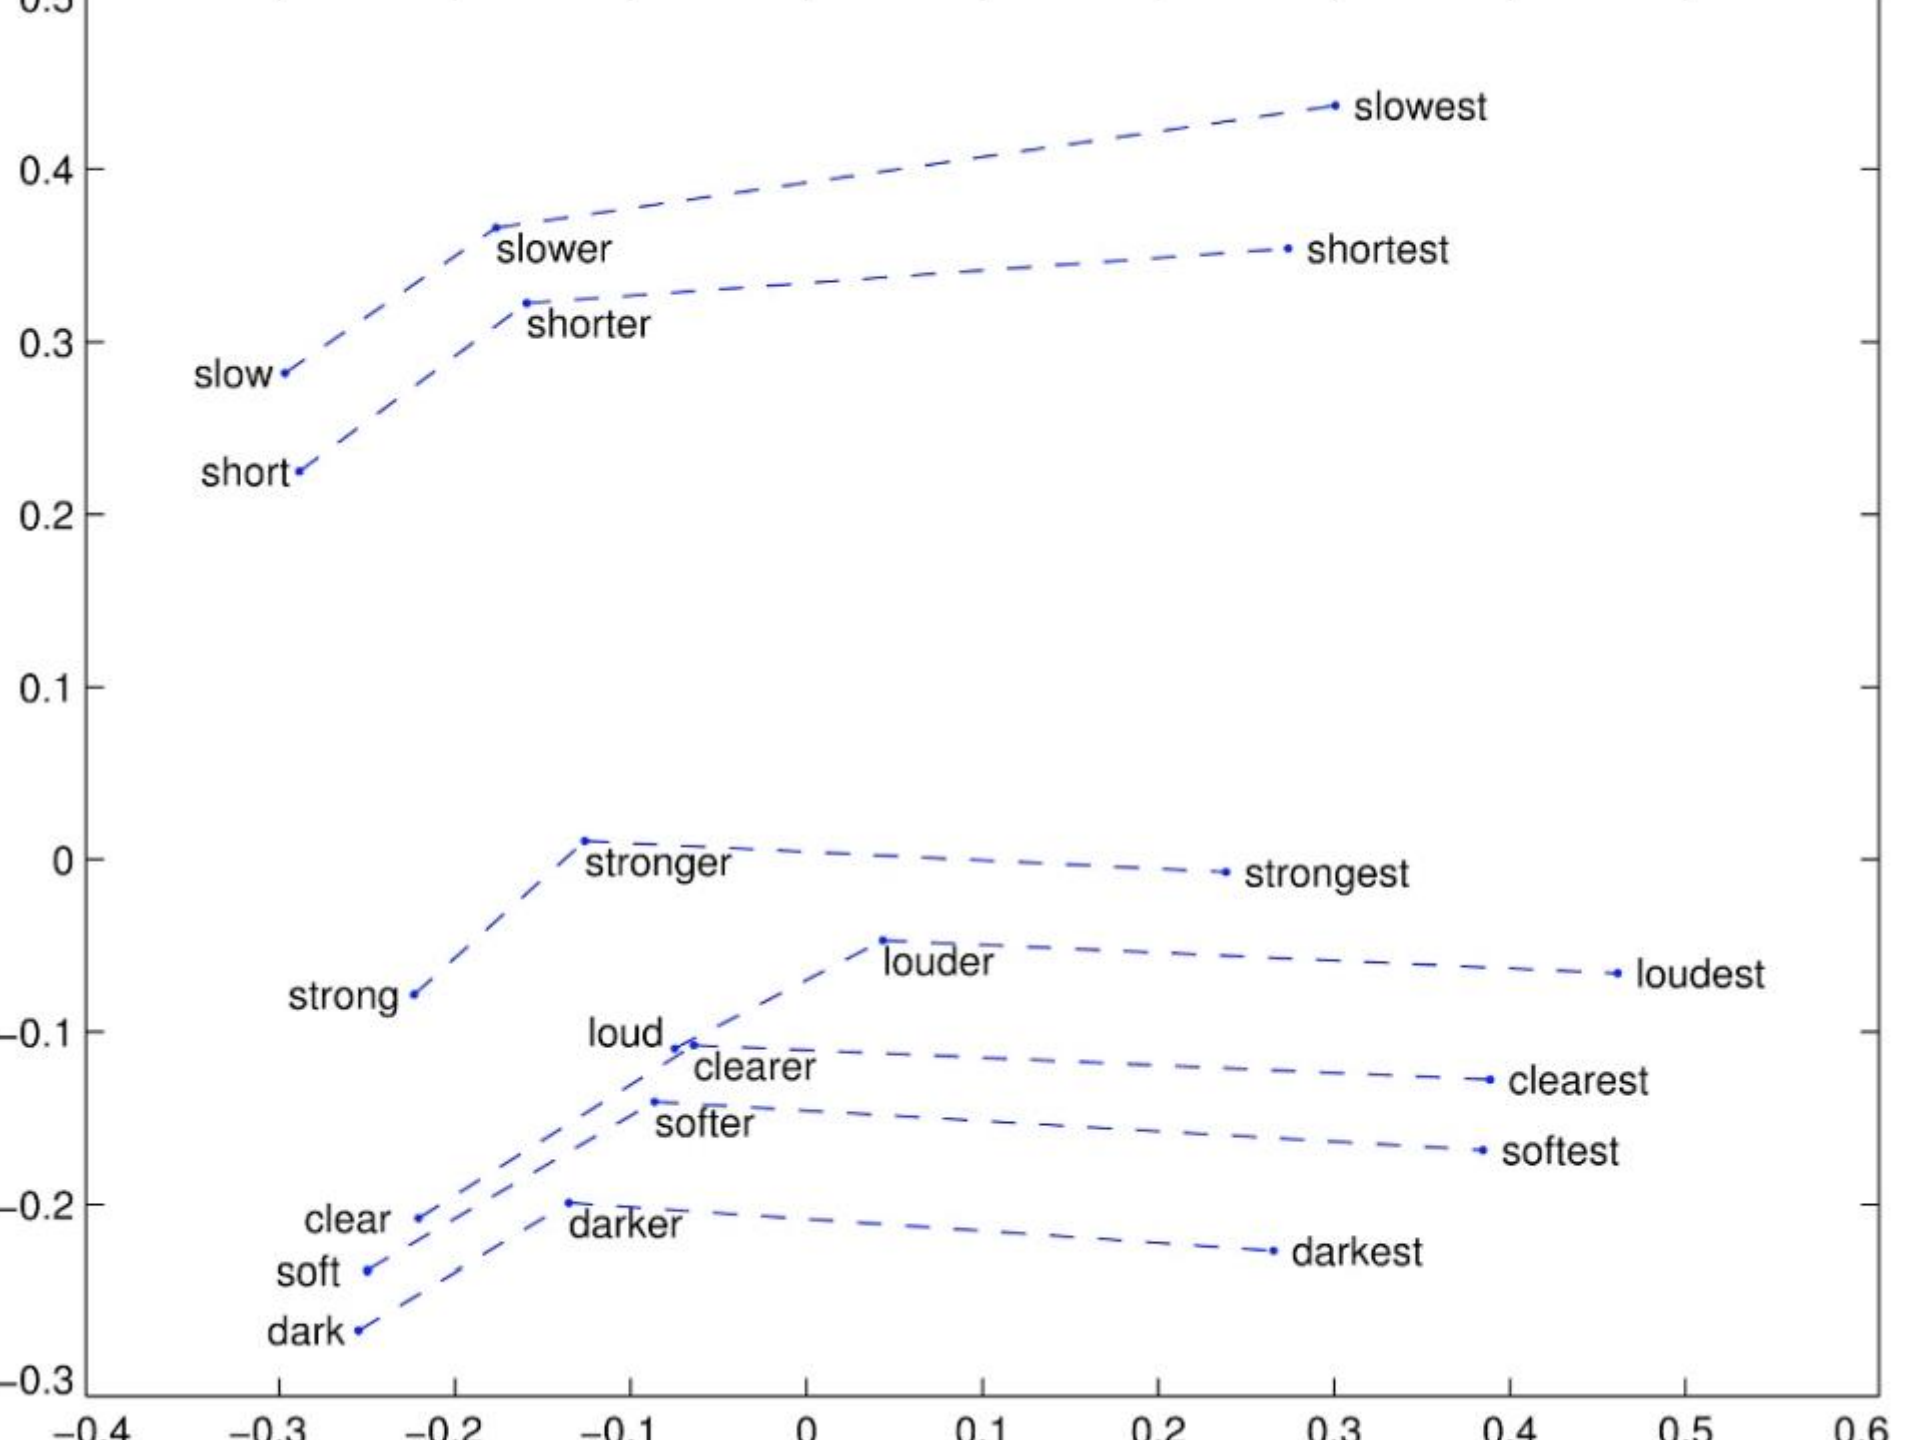
man:woman :: king:?

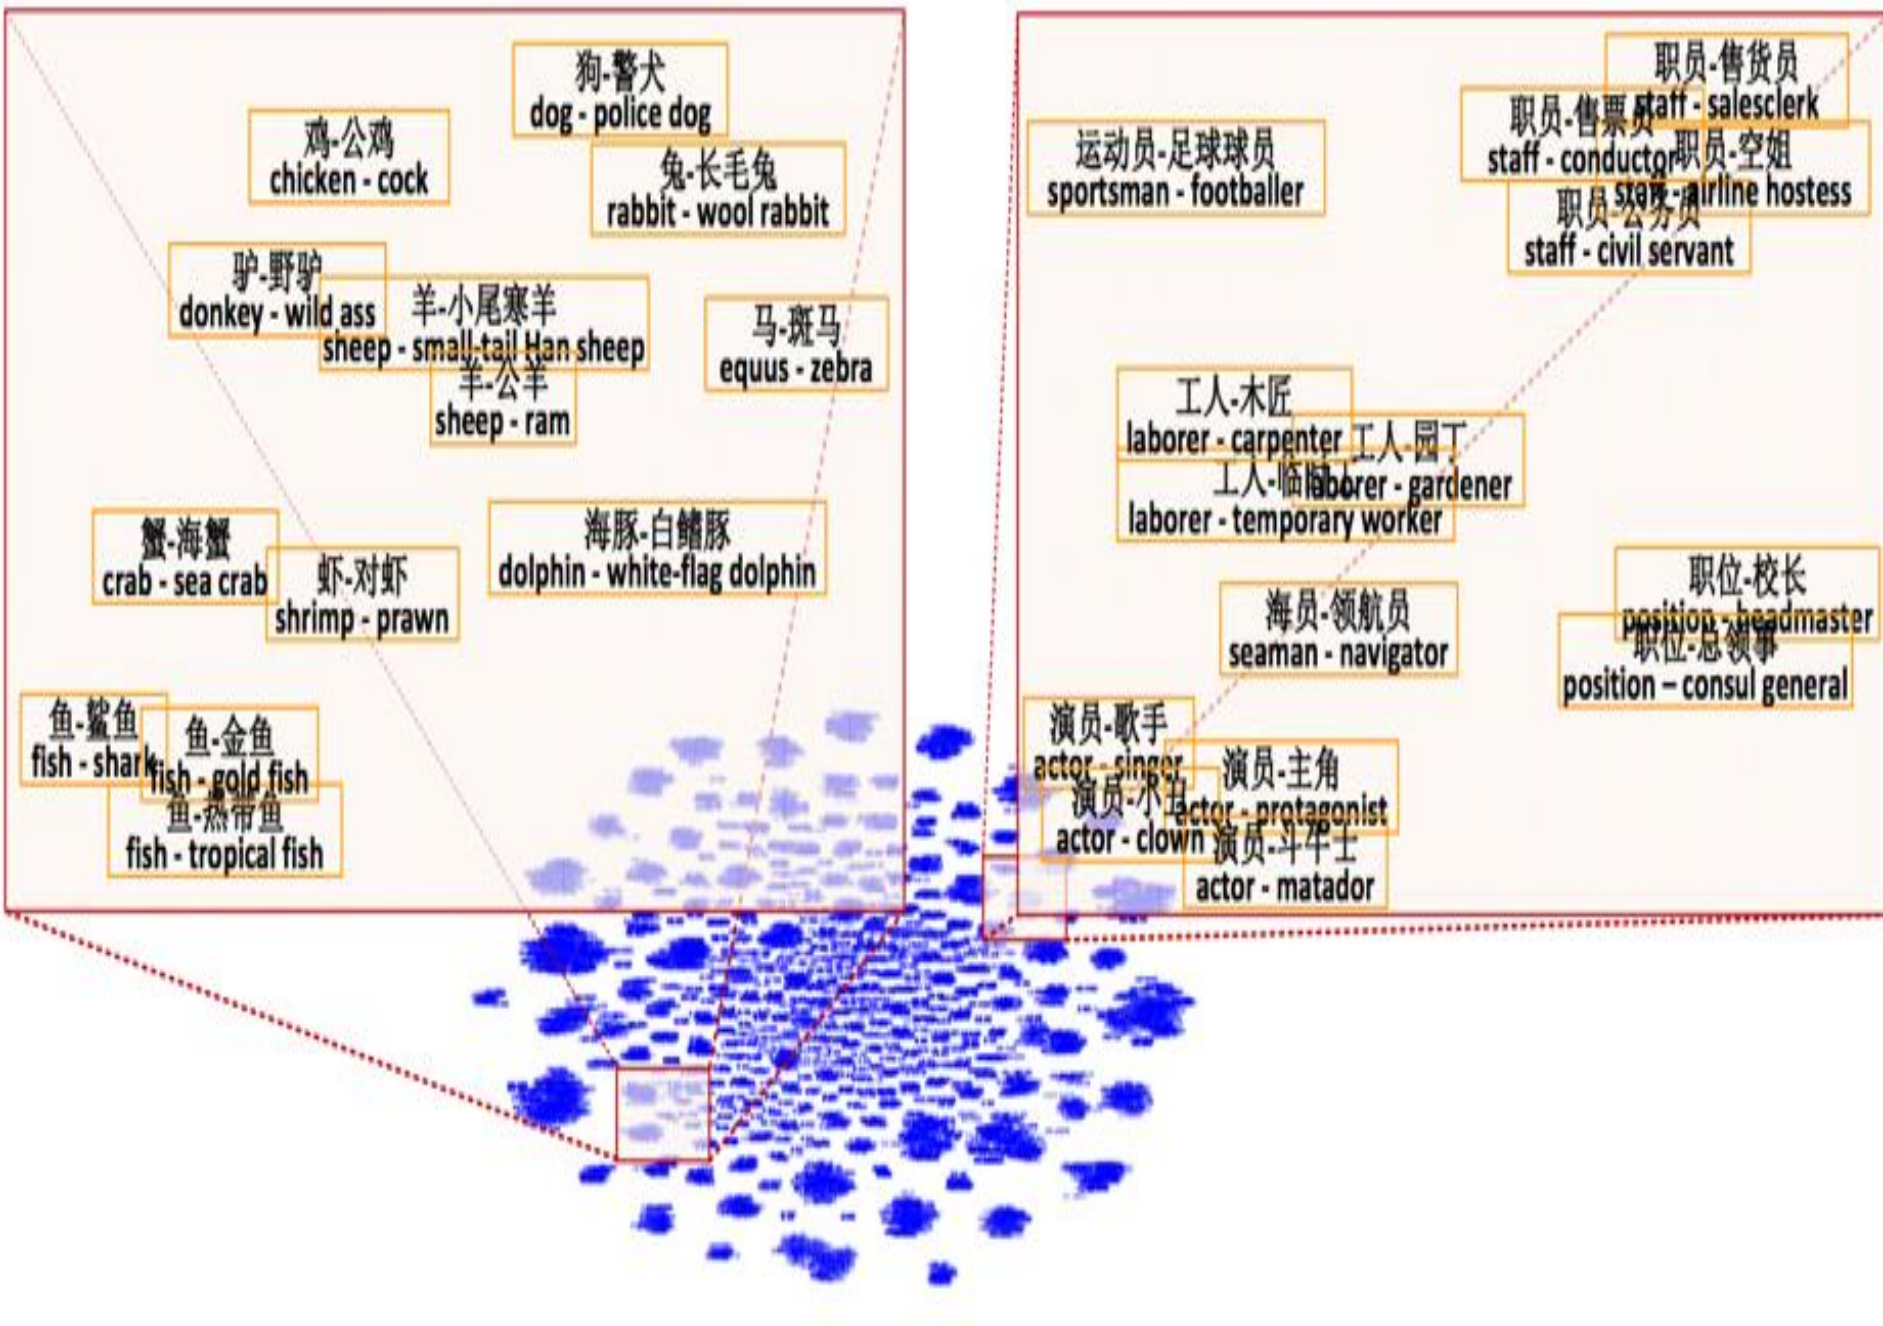
- Evaluate word vectors by how well their cosine distance after addition captures intuitive semantic and syntactic analogy questions
- Discarding the input words from the search!
- Problem: What if the information is there but not linear?







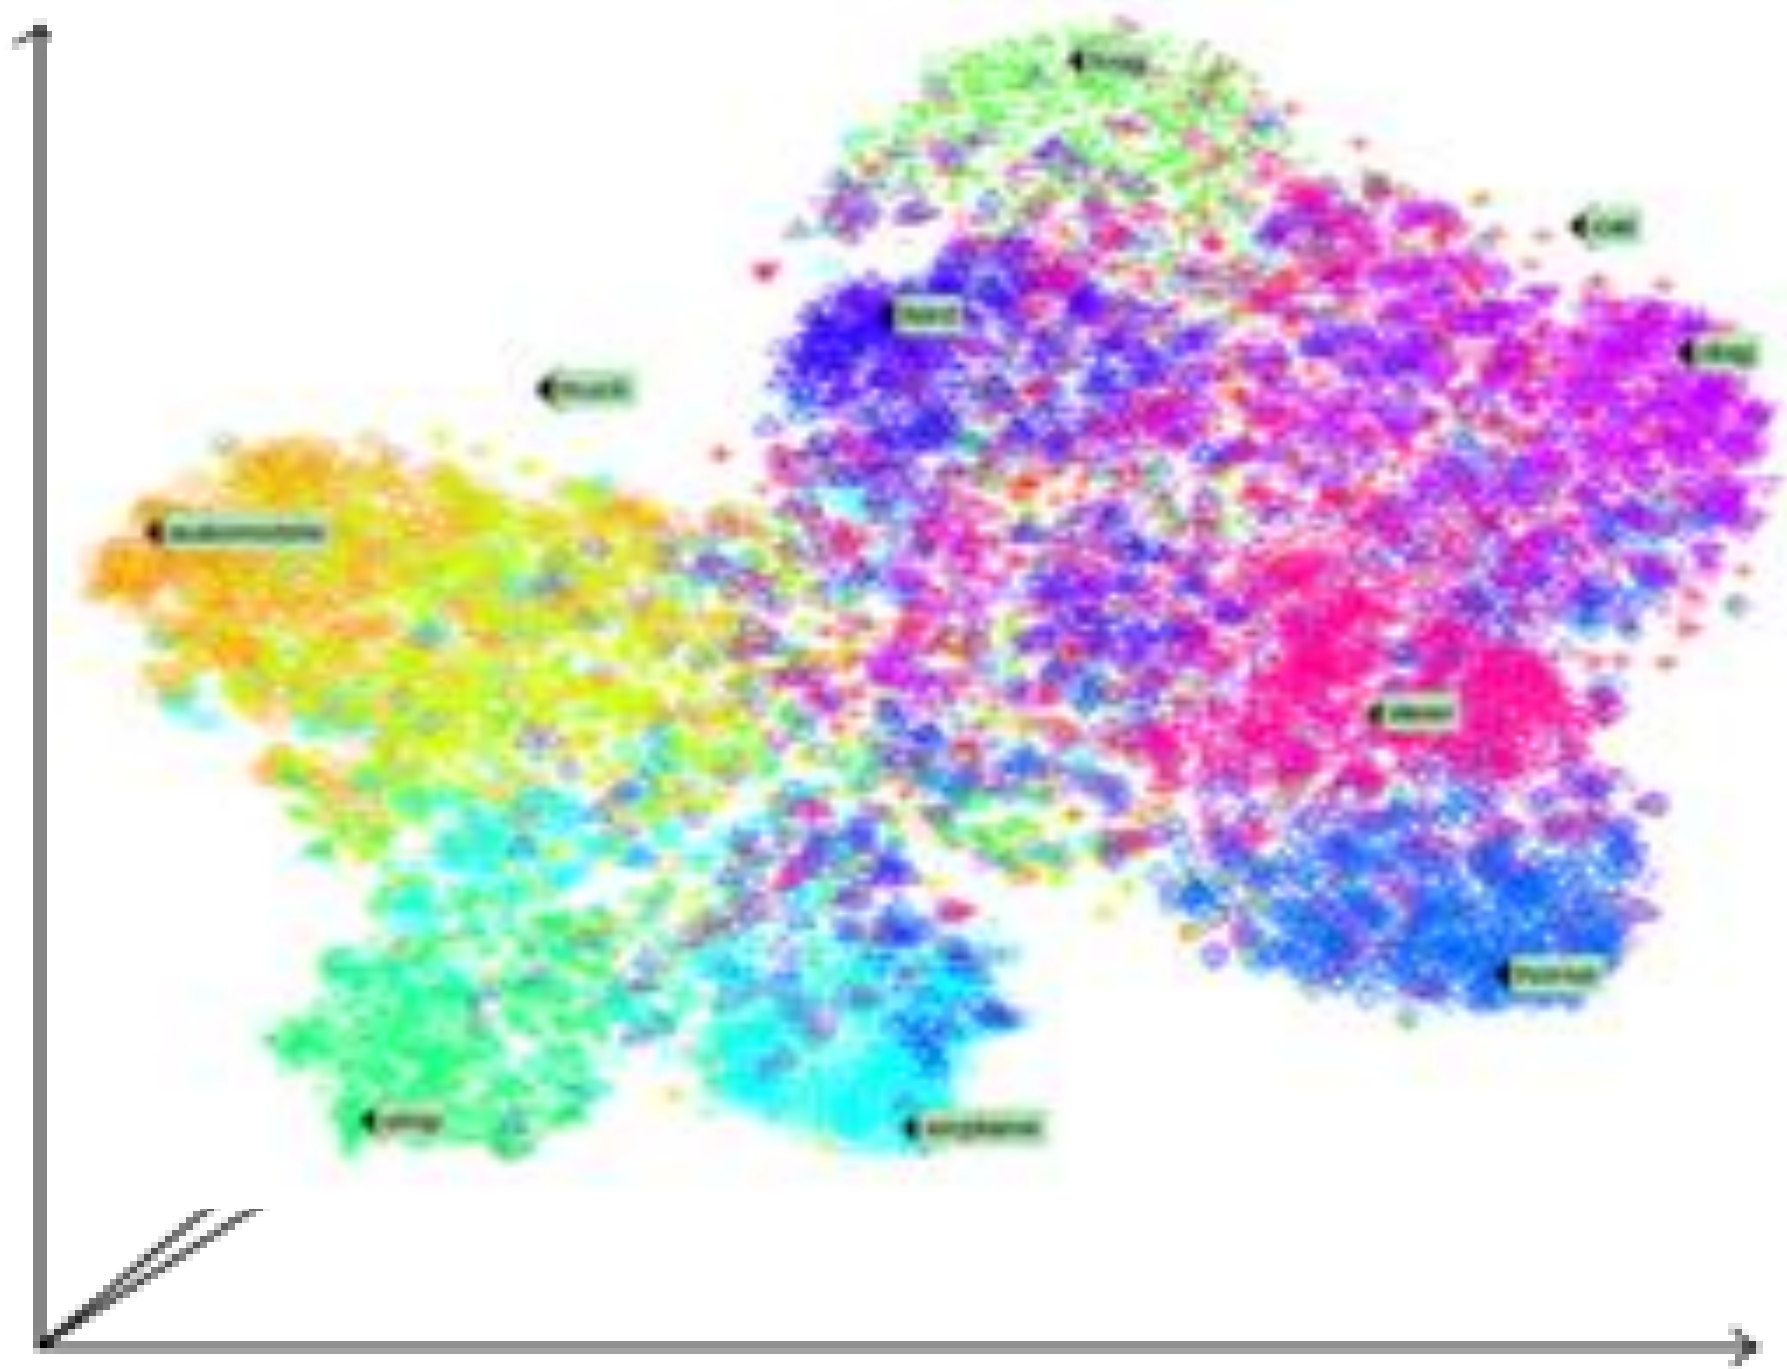




Fu, Ruiji, et al. Learning semantic hierarchies via word embeddings. ACL 2014.

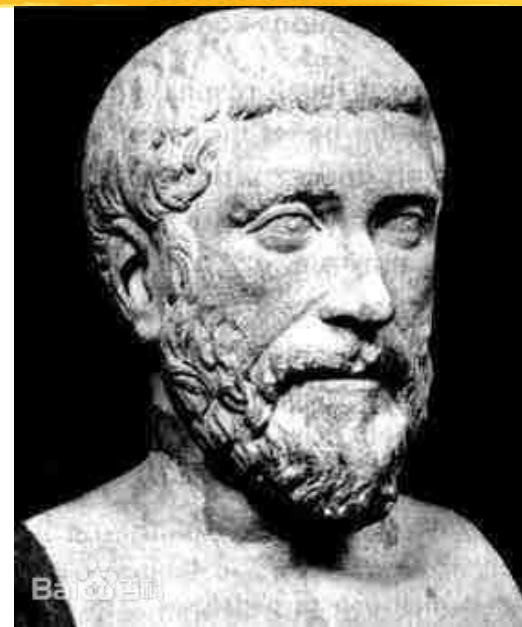






## 3.5 Thesaurus及term自动关联(Word2Vec)

- 毕达哥拉斯：  
“万物皆数”
- 深度学习：  
“万物皆数组”（向量）！



（约前572—约前500）



Yann LeCun



Yoshua Bengio



Geoffrey Hinton

# The skip-gram model with negative sampling (HW2)

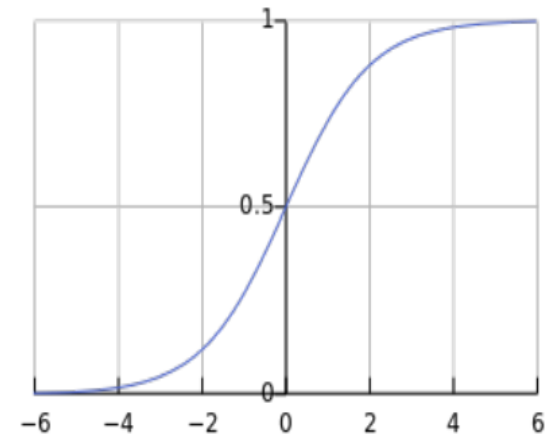
- From paper: “Distributed Representations of Words and Phrases and their Compositionality” (Mikolov et al. 2013)

- Overall objective function (they maximize):

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T J_t(\theta)$$

$$J_t(\theta) = \log \sigma(u_o^T v_c) + \sum_{i=1}^k \mathbb{E}_{j \sim P(w)} [\log \sigma(-u_j^T v_c)]$$

- The logistic/sigmoid function:  $\sigma(x) = \frac{1}{1+e^{-x}}$   
(we'll become good friends soon)
- We maximize the probability of two words co-occurring in first log and minimize probability of noise words





$$\begin{aligned}
 f'(z) &= \left( \frac{1}{1 + e^{-z}} \right)' \\
 &= \frac{e^{-z}}{(1 + e^{-z})^2} \\
 &= \frac{1 + e^{-z} - 1}{(1 + e^{-z})^2} \\
 &= \frac{1}{(1 + e^{-z})} \left( 1 - \frac{1}{(1 + e^{-z})} \right) \\
 &= f(z)(1 - f(z))
 \end{aligned}$$

- Notation more similar to class and HW2:

$$J_{neg-sample}(\mathbf{u}_o, \mathbf{v}_c, U) = -\log \sigma(\mathbf{u}_o^T \mathbf{v}_c) - \sum_{k \in \{K \text{ sampled indices}\}} \log \sigma(-\mathbf{u}_k^T \mathbf{v}_c)$$

- We take  $k$  negative samples (using word probabilities)
- Maximize probability that real outside word appears,  
minimize probability that random words appear around center word
- Sample with  $P(w) = U(w)^{3/4} / Z$ , the unigram distribution  $U(w)$  raised to the 3/4 power  
(We provide this function in the starter code).
- The power makes less frequent words be sampled more often

$$\frac{\partial J(\theta)}{\partial v_c} = -\sigma(-u_o^T v_c)u_o + \sum_{k=1}^K \sigma(u_k^T v_c)u_k$$

$$\frac{\partial J(\theta)}{\partial u_o} = -\sigma(-u_o^T v_c)v_c$$

$$\frac{\partial J(\theta)}{\partial u_k} = \sum_{k=1}^K \sigma(u_k^T v_c)v_c$$

证明说明见：<https://medium.com/analytics-vidhya/maths-behind-word2vec-explained-38d74f32726b>

此外：

$$P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$$